

Small Fish — PMM Foundations: Market, Competitor & Customer Research

2026-09-19 · @Someone

Executive summary

The market for local business lists is proven and crowded at the bottom, and empty at the one layer buyers actually need: relevance. Every tool answers "all dentists in Austin". Almost none answers "dentists in Austin that offer implants and have no online booking", because those facts live on each business's own website, in prose.

The thesis. Small Fish finds local businesses by what they actually do, not by the category Google assigned them. A user types a place, a type of business and a few plain-English criteria. Small Fish returns only the businesses whose own websites prove a match, with the sentence that proves it.

The verdict: go, as a focused indie SaaS product. The business case is a durable \$10K to \$50K MRR product serving people who sell to local businesses, not a venture-scale platform. It is conditional on one test: Small Fish must beat Scrap.io filters and Exa Websets on accuracy for the same niche search.

Five findings that shape everything below:

1. **Demand is proven.** Dozens of paid tools sell local business lists, and the leading filtered tool reports 50,000+ users.
2. **The raw data layer is a commodity.** Google Maps records sell for about \$3 per 1,000; one tool offers 100,000 leads for \$9.90 a month. Nobody wins on data volume.
3. **Every mainstream tool filters on map fields,** not website content. Category, rating, review count and "has website" are the ceiling of today's targeting.
4. **AI qualification is arriving, but for technical users.** Clay recipes, Apify AI actors and Exa Websets can judge relevance, but each needs setup, credits or engineering comfort that the core buyer lacks.
5. **The buyer is anyone who sells to local businesses and is not a GTM engineer:** AI automation agencies, web design and SEO shops, local marketing agencies, and SaaS companies selling to small businesses. They pay \$10 to \$100 a month and buy tools rather than build them.

Market definition and sizing

The US market for local business list tools is roughly \$150M to \$250M a year, and a focused product can realistically capture \$400K to \$600K ARR within three years. The target supply is huge; the buyer pool is large, fast-growing and price-sensitive.

The job

Help me find the local businesses worth selling to, and tell me why. The category today is called Google Maps scrapers or local lead generation. The job buyers actually hire it for is narrower: a short, relevant list they can contact with confidence.

The accounts being searched (supply)

Measure	Count	Source
US small businesses (under 500 staff)	36.2M	SBA via uscompanieslist
Employer firms (1+ employee)	5.5M to 6.4M, depending on source	Census
Map-visible US businesses	18.4M	CUFinder
Map-visible with zero reviews	8.2M (44.5%)	CUFinder
New business applications, Jan 2026	532,319 in one month	Wave Connect citing Census

Two implications. First, the addressable universe for Small Fish is the map-visible businesses that also have a readable website, likely a few million, not 36M. Second, churn is extreme: roughly half of new employer businesses close within five years, so every static list decays. Freshness is a real product feature, not a nice-to-have.

The people buying (demand)

Buyer group	US count	Signal	Source
Digital advertising agencies	100,202	+14.9% in 2026; 16.6% CAGR 2021-26	IBISWorld

Buyer group	US count	Signal	Source
All advertising agencies	114,014	+4.4% in 2026	IBISWorld
AI automation agencies	Not counted by any bureau	Fastest-growing new agency type; AI receptionist market growing ~24% a year	SchedulingKit
Web design and SEO shops	Tens of thousands (estimate)	Core buyers of Maps scrapers today	Open question
SaaS and B2B services selling to SMBs	Thousands of companies (estimate)	Outbound teams with the highest willingness to pay	Open question

Sizing

Level	Definition	Estimate	Basis
TAM	US sellers to local businesses who could use a list tool	\$150M to \$250M a year	~250K to 400K buyers × ~\$600 a year
SAM	Those who prospect outbound and already pay for a list or lead tool	\$50M to \$80M a year	~30% of TAM, consistent with 50,000+ users on one mid-tier tool
SOM (3 years)	Realistic Small Fish share	\$400K to \$600K ARR	700 to 1,000 paying users at ~\$50 a month

English-speaking markets outside the US (UK, Canada, Australia) add an estimated 30% to 40% on top. All buyer-side figures are estimates built from agency counts; treat them as orders of magnitude, not forecasts.

Why now

- **Reading websites became cheap.** Small, fast models can judge a homepage for a fraction of a cent, which makes website-level filtering affordable at list-tool prices.
- **A new buyer class appeared.** AI automation agencies sell voice receptionists, booking and follow-up bots to clinics and trades, and every one of them needs to find businesses with a specific gap, such as no online booking.

- **The old tools are stuck on map fields.** Their data model predates cheap reading, so relevance is structurally hard for them to add well.

Competitive landscape

Nobody owns local relevance yet. The category leader, Scrap.io, wins on volume and map-data filters; Exa Websets wins on AI verification but is generic and caps results; a long tail of small AI tools has tried the exact idea and not executed it. The opening is quality and trust, not the concept.

Five tiers

Tier	Players	What they sell	Ceiling
1. Raw extractors	Outscraper, Apify, Leads Sniper, Map Lead Scraper	Cheap rows from Google Maps, pay per record	No judgment; user cleans everything
2. List tools	D7 Lead Finder, LeadScrape, LeadSwift	Keyword + city → CSV	Static data, 4 basic filters
3. Filtered list platforms	Scrap.io	20+ map and tech filters, real-time extraction, API, MCP	Filters on listing and tech fields, not what the business does
4. AI list builders	Exa Websets, Clay, Apify AI actors	Natural-language criteria, verification per row	Generic, technical, or credit-heavy; not local-native
5. AI-native local tools	LocalLead.ai, B2BLeadFinder, LitarelAI	"Describe your ideal client" for local businesses	Thin products; mostly "digital gap" pitches for web agencies

Pricing snapshot, September 2026

Product	Entry price	What entry buys	Billing unit
Outscraper	Pay as you go	~\$3 per 1,000 records; 500 free a month	Per record

Product	Entry price	What entry buys	Billing unit
Map Lead Scraper	\$9.90/month	100,000 leads	Flat
Scrap.io	\$35/month (annual)	10,000 exports; 7-day trial with 100 leads	Per export
D7 Lead Finder	\$44.99/month	~2,500 searches; no trial	Per search
LeadScrape	\$97-247/year	Unlimited searches, desktop	Flat licence
Exa Websets	\$49/month (Core)	8,000 credits ≈ 800 verified results	10 credits per full match; non-matches free
Clay	\$185/month (Launch)	2,500 data credits, 15,000 actions	Two meters
Apify AI lead qualifier	Pay as you go	\$5 per 1,000 with AI scoring	Per result

The floor is brutal. A raw row costs a third of a cent; an AI-scored row costs half a cent on Apify. Exa's verified result costs about 6 cents on Core. Small Fish must price against the value of a *match*, not the cost of a row.

Profiles of the three that matter

Scrap.io: the incumbent to beat. Real-time extraction from Google Maps across 195 countries, 225M+ establishments indexed, 50,000+ users, G2 4.9. Filters include website and email presence, social profiles, claimed listing, rating and review ranges, contact form on site, ad pixels, first-detection date and duplicate exclusion. It ships an API, a Make module and an MCP server. It markets hard with comparison pages against every rival. [Source](#).

Where it stops: every filter describes the listing or the site's technology, never what the business does. It can find dentists with a contact form and a Meta pixel. It cannot find dentists that offer implants, or salons that do bridal packages, or HVAC firms that do commercial work. It also sells volume (10,000 exports is the entry plan), which is the opposite of a short, trusted list.

Exa Websets: the closest technical match. Plain-English criteria, agentic search, and a verification status with reasoning and references for every result. Non-matches cost nothing. Exa raised a \$250M Series C at a reported \$2.2B valuation in May 2026, so it will keep improving. [Source](#), [source](#).

Where it stops: it is a general web index built for knowledge workers across people, companies and papers, not a local-business product. Core caps each Webset at 100 results, large searches can take up to an hour, and there is no map-native geography, no local category model, and no outreach-ready local fields. Its coverage of tiny local businesses with thin sites is unproven and must be tested.

Clay: the builder's option. Publishes a full recipe for scraping Google Maps by category and geography, finding owners through a waterfall across ~50 sources, and letting AI read sites. [Source](#).

Where it stops: the buyer must build and maintain the table, and pay two meters. Clay's own local guide ranks leads by rating and review count, not by relevance to what the seller offers.

The long tail tells us the idea is obvious and the execution is not

- **LocalLead.ai** already pitches "AI converts your requirements into searchable keywords... analyses each lead against your requirements with scores and reasons" plus monthly fresh batches. Its site footer reads 2023 and its testimonial uses a stock template avatar: an idea without traction. [Source](#).
- **B2BLeadFinder** and a cluster of Apify actors sell "digital gaps": no website, low reviews, missing hours, no booking. These are proxy signals for web and SEO agencies, not relevance to an arbitrary offer. [Source](#), [source](#).
- **LitarelAI** bundles natural-language lead discovery into a local SEO suite. [Source](#).

The lesson: "describe your ideal client and AI scores leads" is not a moat or even a differentiator on its own. Winning requires accuracy that is visibly better, trust built on shown evidence, and distribution.

Capability matrix

Capability	Raw extractors	Scrap.io	Exa Websets	Clay	AI-native local tools	Small Fish (target)
Map-native local coverage	Yes	Yes	Weak	Via recipe	Yes	Yes

Capability	Raw extractors	Scrap.io	Exa Websets	Clay	AI-native local tools	Small Fish (target)
Filter on listing fields	Some	20+	No	Yes	Some	Yes
Filter on what the business does (site content)	No	No	Yes	If built	Claimed	Core
Evidence sentence per match	No	No	Yes	If built	Partial	Core
Honest "couldn't tell"	No	No	Partial	If built	No	Core
Pay only for matches	No	No	Yes	No	Unclear	Yes
No setup for a non-technical user	Partial	Yes	Yes	No	Yes	Yes
Saved searches with new-match alerts	No	No	Monitors (paid)	Build it	Monthly batch (claimed)	Yes

What users complain about today

- **Stale data.** The most repeated complaint about D7 in its own reviews; D7-sourced emails reportedly run ~9% invalid versus ~2% from live extraction. [Source](#).
- **Commitment before proof.** "I needed only few lists so buying a subscription is not rational to me", quoted from r/LeadGeneration. [Source](#).
- **Cleanup overhead.** Raw scrapers leave all filtering to spreadsheets, and reviewers note the cleanup eats the savings on steady prospecting. [Source](#).
- **Unpredictable AI costs.** Exa's agent runs range from \$0.012 to \$1.00 each depending on effort, which reviewers call hard to forecast. [Source](#).

Customer research

The best first customer is the AI automation agency owner: new, growing, non-technical about data, and selling an offer that only fits businesses with a specific, website-visible gap. Web design and SEO agencies are the largest volume segment; SaaS teams selling to small businesses pay the most. All three share one pain: map categories cannot describe what a business actually does.

The structural root cause

Google's own rules make relevance impossible to filter on Maps. Categories must complete the sentence "this business **is** a", not "this business **has** a", and must not be used to describe services, products or attributes. [Google Business Profile guidelines](#). So "offers implants", "does commercial work" or "takes online bookings" can never be a map filter. It lives on the website.

The website is also more current than the listing. In a sample of 230 local businesses, 53% had at least one mismatch between their Google profile and their website. [DreamHost](#).

Segments

Segment	Size signal	What they sell to local businesses	Criteria they need	Willingness to pay	Priority
AI automation agencies	Fast-growing; AI receptionist market ~24% a year; 72% of deployments in healthcare, legal, home services (SchedulingKit)	Voice receptionists, booking bots, follow-up automation	No online booking, phone-only intake, after-hours gaps, high call volume	\$29-99/month; one client pays \$500-15,000	1
Web design, SEO and local marketing agencies	~100,000 digital agencies, +14.9% in 2026 (IBISWorld)	Sites, SEO, ads, reputation	Outdated site, no booking, specific services offered, no blog, weak reviews	\$29-79/month	2
SaaS and B2B services selling to small businesses	Thousands of vendors: practice software, POS, payroll, insurance, IT	Software and services for one vertical	Offers X service, uses Y tool, size band, independent	\$79-299/month	3
Freelance list builders	Upwork list-building rates \$13-45/hour	Lists to their own clients	Whatever the client asks	\$49-149/month; resells output	4 (power users)

Personas

Arjun, the AI automation agency founder (primary). One to three people, 6 to 18 months in business, sells an AI receptionist and booking automation to dental and med spa clinics at \$300 to \$1,000 a month per client. Learned the business from YouTube and Skool

communities. Uses n8n or Make, Instantly for email, a Maps scraper for lists.

- **Job:** "Find me clinics that are clearly losing bookings, so my pitch lands."
- **Today:** scrapes 2,000 clinics, opens websites one by one to check for online booking, gives up after 40.
- **Pain:** 80% of his list already has booking software; his emails read as irrelevant; replies are under 1%.
- **Win:** "212 med spas in Dallas with no online booking, and the page that proves it."

Maria, the local web and SEO agency owner (volume). Five-person shop, serves home-service trades, sells \$2,000 to \$8,000 websites and \$500 to \$1,500 monthly SEO.

- **Job:** "Find trades whose website is holding them back."
- **Today:** uses Scrap.io or D7 and filters "has website" and review count.
- **Pain:** can filter by review count but not by "site looks 2014, no service pages, no quote form".
- **Win:** a list of HVAC firms that do commercial work, with no quote form and no service-area pages.

Dev, the SaaS go-to-market lead (highest value). Sells scheduling or practice software to one vertical, runs a two-person outbound team.

- **Job:** "Every independent practice in my territory that fits our product, and not the ones that already use a competitor."
- **Today:** Apollo covers the vertical poorly; interns build lists by hand.
- **Pain:** cannot filter by "independent, not part of a group" or "uses competitor X".
- **Win:** a clean, evidence-backed territory list, refreshed when new practices open.

Jobs to be done

When I...	I want to...	So I can...
Start prospecting a new niche or city	Get only the businesses my offer fits	Stop wasting sends and credits on bad fits
Write a cold email	Know one true, specific thing about each business	Sound researched, not spammy
Pitch a client on a gap	Show proof of that gap	Win the argument before the call

When I...	I want to...	So I can...
Hand a list to my team or client	Trust it without re-checking	Save hours and my reputation
Come back next month	See only what's new	Keep prospecting without re-buying the same list

Pains, ranked

1. **Irrelevant lists.** Category search returns every business of a type; most do not have the gap the seller fixes.
2. **Manual checking.** Qualifying means opening websites one by one, which does not scale past a few dozen.
3. **Paying for junk.** Almost every tool charges per record exported, including the 70% to 90% that are wrong fits.
4. **Generic outreach.** Without a specific fact, emails look templated and replies collapse.
5. **Stale data and bounces.** Static databases decay; businesses close or change fast.
6. **Unpredictable AI costs.** Credit-based AI tools make a budget hard to set.

How they buy

- **Discovery:** YouTube tutorials, Skool and Reddit communities, "best Google Maps scraper" comparison pages, AppSumo.
- **Evaluation:** a free trial or free count first. Buyers resent subscriptions before proof; one quoted user says buying a subscription for a few lists "is not rational to me".
[Source.](#)
- **Decision maker:** the founder, alone, on a card, in one session.
- **Price anchors in their head:** Scrap.io \$35–139/month, Exa \$49/month, freelancers at \$13–45/hour, a single client worth \$500 to \$15,000.
- **Switching trigger:** a bad campaign on an irrelevant list, or a new offer that needs a new niche.

Evidence gaps to close with interviews

The segment and pain picture above is built from public sources, community content and competitor positioning, not yet from first-hand interviews. Before the product document is final, run 10 interviews: 5 AI automation agency owners, 3 web and SEO agencies, 2 SaaS

go-to-market leads. Test three things: how they qualify today, what they would pay per matched lead, and which criteria they would type first.

Gap analysis

The white space is a local-native tool that filters on what businesses actually do, proves every match, and charges only for matches, built for a non-technical seller. Each piece exists somewhere; no product combines them.

Where each player sits

Two axes explain the market: how local-native the tool is (map geography, local categories, local contact fields) and how much it judges relevance from the business's own website.

	Low relevance judgment	High relevance judgment
Local-native	Scrap.io, D7, Outscraper, LeadScrape, LeadSwift	Empty, except thin attempts (LocalLead.ai, B2BLeadFinder)
Generic web	Apollo, database vendors	Exa Websets, Clay (if built)

The gaps, ranked by value to the buyer

#	Gap	Who has part of it	Why it's still open	Size of prize
1	Filter on services, offers and operating model stated on the site	Exa (generic), Clay (build it)	Incumbents are built on map fields; Google rules forbid services as categories	Core of the product
2	Evidence for every match, shown next to the row	Exa	Scrapers never read the site in depth; showing proof needs a different UI	The trust layer; drives conversion from trial
3	Honest "couldn't tell" instead of a forced yes or no	Nobody well	Tools fear looking incomplete	Builds credibility; cuts bounces and bad sends

#	Gap	Who has part of it	Why it's still open	Size of prize
4	Pay only for matches	Exa	Scrapers' economics depend on selling volume	Removes the "paying for junk" pain directly
5	A free match count before paying	Scrap.io offers free counts, not match counts	Needs relevance judgment up front	The acquisition hook
6	Saved searches that alert on new matches	Exa Monitors (paid add-on), LocalLead (claimed)	Requires diffing and re-reading	Retention engine
7	Criteria templates by offer ("AI receptionist prospects")	Nobody	Requires understanding sellers' offers, not just data	Faster time to first value; SEO surface

Why the space is still empty

- **Incumbent economics.** Scrap.io and the raw extractors sell volume. A product that returns 212 matches instead of 2,000 rows threatens their pricing model.
- **Cost until recently.** Reading every website with a capable model was too expensive at \$3-per-1,000 price points. Small, fast models changed that in 2025–26.
- **Exa is aimed elsewhere.** A \$2.2B infrastructure company optimises for developers and knowledge workers across the whole web, not for a dental-clinic list at \$29 a month.
- **Small AI entrants lack depth.** The existing attempts stop at "digital gaps" or a marketing claim; none shows evidence, handles uncertainty or proves accuracy.

What would close the gap, and how fast

Threat	Likelihood in 12 months	Impact	Our answer
Scrap.io adds "AI criteria" to its filters	Medium	High	Accuracy benchmarks published openly; evidence UX; pay-per-match pricing they can't copy without hurting revenue
Exa launches a local mode	Low to medium	High	Offer templates, local fields, and price and UX for non-technical sellers
A Clay template goes viral	Medium	Medium	Clay's buyer builds; ours won't
New Apify or Gumroad clones	High	Low	Commodity; they lack trust, accuracy and support

Positioning foundations

Small Fish is the local lead finder that searches by what businesses actually do, and proves every match. That single idea should carry the product, the pricing and every page of marketing.

Positioning statement

For agencies and companies that sell to local businesses, **who** waste hours and sends on lists that don't fit their offer, **Small Fish** is a local lead finder that searches businesses by what their own websites say they do. **Unlike** Google Maps scrapers, which filter only on category, ratings and listing fields, Small Fish reads every business's site, returns only real matches with the sentence that proves each one, and charges only for matches.

Category

Use the category buyers already search for, **local lead finder** (with "Google Maps scraper alternative" for search), and own a modifier: **relevance-first**. Do not invent a new category name; there is no budget to teach one.

Value proposition

- **Headline:** Find local businesses by what they actually do.
- **Sub-head:** Type a place, a type of business, and what you're looking for. Get only the businesses whose own websites prove a match, and pay only for those.
- **Tagline:** The big nets miss the small fish.

Three pillars

Pillar	Claim	Proof to build
Relevance	Search by services, offers, gaps and operating model, not just category	Side-by-side: Scrap.io category search vs Small Fish criteria search on the same city
Proof	Every match shows the sentence and page that prove it; unclear cases are marked, not guessed	Published accuracy benchmark on 3 niches, checked by hand
Fair price	Free match count up front; pay only for matches	Pricing page showing cost per match vs cost per useful row elsewhere

Differentiators, and what they are not

- **Is:** website-based relevance, evidence per row, honest uncertainty, pay-per-match, offer templates, alerts for new matches.
- **Is not:** the biggest database, the cheapest per row, an email sender, a CRM, or an enrichment workbench. Say so plainly; it builds trust with buyers burned by overpromising tools.

Messaging by segment

Segment	Hook	Example search to demo
AI automation agencies	"Find clinics that are losing bookings, before you pitch them."	Med spas in Dallas with no online booking

Segment	Hook	Example search to demo
Web and SEO agencies	"Find the businesses your service actually fixes."	HVAC firms in Ohio that do commercial work and have no quote form
SaaS selling to small businesses	"Your whole market, qualified by what they do, not what Apollo guessed."	Independent dental practices in Texas offering implants, not part of a group
Freelance list builders	"Deliver lists your clients never have to re-check."	Any client brief, with proof columns

Proof points to create before launch

1. **The accuracy benchmark.** Three niches, 100 businesses each, hand-checked. Publish precision and the "couldn't tell" rate.
2. **The side-by-side.** Same city and niche in Scrap.io, Exa Websets and Small Fish: rows returned, true matches, cost per true match.
3. **One customer story** with a reply-rate change from a relevance-filtered list.

Risks, constraints and the data landscape

The biggest structural constraint is data rights: Google Places content can't be stored, so Small Fish must build its own record from open data and each business's own website. The biggest market risk is an incumbent adding AI criteria; the biggest product risk is accuracy that isn't visibly better.

Candidate data sources

Source	Scale	Cost	Can we store it?	Catch
Google Places API (New)	Best US coverage	\$32–35 per 1,000 for useful Text Search fields	Only place IDs indefinitely; coordinates 30 days; names, phones, hours must be re-fetched (source)	Max 60 results per query; renting the same rows repeatedly

Source	Scale	Cost	Can we store it?	Catch
Overture Maps Places	72M+ places worldwide (Jan 2026)	Free download	Yes, per-source licences (CDLA, Apache, CC0)	Needs engineering to query, dedupe and validate; fill rates vary
Foursquare OS Places	106M+ places (Dec 2025)	Free	Yes, Apache 2.0	Same as above; partly merged into Overture
Third-party Maps extractors (Outscraper, Apify)	Full Maps coverage	~\$3 per 1,000	Contractual grey zone	Relies on scraping Google Maps, which Google's terms prohibit
Public registries (licensing boards, state business filings)	Vertical-specific	Mostly free	Yes	Different format per state and vertical
The business's own website	Every business with a site	Fetch cost only	Our derived judgment, yes; store extracted facts, not full copies	Must respect robots.txt and rate limits

Open data is bigger but messier: one 2026 benchmark found OpenStreetMap covered only 26% of verified locations for one chain in a test market and filled 39.8% of attributes. [Source](#). Overture and Foursquare are far larger than OSM, but coverage of small local businesses must be measured per niche before launch.

Recommended data posture: candidates from Overture and Foursquare open data first; Google Places API only to fill gaps, storing place IDs; the business's website as the source of truth for every judgment; no direct Google Maps scraping.

Legal and compliance

- **Scraping Google Maps directly** breaks Google's terms of service. Court rulings such as *hiQ v. LinkedIn* concern federal computer-crime law, not contract terms, so they do not make it safe. Avoid it.

- **Website reading** of public business pages is standard practice. Honour robots.txt, identify the crawler, throttle requests, and store extracted facts rather than full page copies.
- **Outreach law is the user's responsibility** (CAN-SPAM in the US, GDPR and PECR in the UK and EU), but the product should make compliance easy: no invented emails, source shown for every contact, opt-out-friendly exports.
- **Personal data:** keep to business information that businesses publish about themselves; avoid scraping personal social profiles or review authors.

Risk register

Risk	Likelihood	Impact	Mitigation
Accuracy not clearly better than Exa or Scrap.io filters	Medium	Fatal	Run the benchmark before building; publish it; build only if Small Fish wins
Scrap.io adds AI criteria	Medium	High	Evidence UX, honest uncertainty and pay-per-match are hard for a volume seller to copy
Open data misses too many small businesses	Medium	High	Measure coverage per niche; fill gaps with Places API IDs
Model costs erode margin at low prices	Low	Medium	Small models; read homepage plus one page; cache judgments; price per match
Low-price buyers churn fast	High	Medium	Saved searches with new-match alerts; annual plans; agency tier
Users chase volume, not relevance	Medium	Medium	Position explicitly for quality; don't compete on row counts
AI agents do the job for free	Medium	Medium	Local coverage, speed, saved searches and exports that a chat window lacks

Implications for product and GTM

The research points to a narrow, sharp product: relevance-first search for local businesses, sold self-serve to people who sell to them. Everything below feeds the product document and the GTM plan.

For the product document

1. **Relevance is the product.** Plain-English criteria judged against each business's website is the core loop; map filters are table stakes, not the pitch.
2. **Evidence on every row.** The proof sentence and source page are part of the result, not a hidden detail.
3. **Three answers, not two:** match, no match, couldn't tell. Uncertainty is shown, never forced.
4. **Free match count before payment.** The count is the hook and the demo.
5. **Pay per match.** Non-matches and "couldn't tell" are free.
6. **Offer templates.** Ready-made criteria for common sellers, such as AI receptionist prospects and web redesign prospects.
7. **Saved searches with new-match alerts,** as the retention engine.
8. **Own data layer:** open places data plus website reading, with Google place IDs only as a gap-filler.

For the GTM plan

1. **First segment:** AI automation agencies selling to clinics and trades; second, web and SEO agencies.
2. **Channels:** comparison pages against Scrap.io, D7 and Outscraper; programmatic niche-and-city pages; YouTube and agency communities; cold email using Small Fish's own lists.
3. **Launch proof:** the published accuracy benchmark and the side-by-side comparison.
4. **Price inside the anchors** buyers already hold: roughly \$29 to \$99 a month, per match.

Open questions

- ☐ Does Small Fish beat Exa Websets and Scrap.io on accuracy for three real niche searches?
- ☐ How much of each target niche do Overture and Foursquare cover versus Google Maps?

- ☐ What share of local businesses have a website with enough text to judge?
- ☐ What do buyers pay per matched lead? (10 interviews)
- ☐ Which three criteria do AI agencies type first?
- ☐ Model cost per business read at target accuracy?

Sources

- [Census: Small Business Week 2026](#)
- [US business counts by state, SBA 2026](#)
- [CUFinder: map-visible US businesses](#)
- [Wave Connect: small business statistics 2026](#)
- [IBISWorld: digital advertising agencies, number of businesses](#)
- [IBISWorld: advertising agencies, number of businesses](#)
- [SchedulingKit: AI receptionist statistics](#)
- [Scrap.io vs D7 comparison](#)
- [GBP Leads: Scrap.io alternatives](#)
- [WebLeads: D7 alternatives](#)
- [LeadScrape: lead generation tools compared](#)
- [Exa Websets FAQ and pricing](#)
- [UsagePricing: Exa funding and pricing history](#)
- [TinyFish: Exa pricing](#)
- [Autotouch: Exa vs Clay pricing](#)
- [Clay: scrape Google Maps for local leads](#)
- [Apify: local lead qualifier](#)
- [Apify: no-website lead finder](#)
- [LocalLead.ai](#)
- [B2BLeadFinder](#)
- [LitareIAI](#)
- [Google Business Profile guidelines](#)
- [DreamHost: GBP and website mismatches](#)
- [Open Places API: Google Places pricing](#)
- [CrustAPI: Google Places caching terms](#)
- [Overture Maps: Places guide](#)

- [SafeGraph: open POI data benchmark](#)