

Small Fish — Product Document

2026-09-19 · @Someone

Summary

Small Fish finds local businesses by what they actually do, proves every match, and charges only for matches.

A user types a place, a type of business and what they are looking for, in plain English: "med spas in Dallas that offer Botox and have no online booking." Small Fish reads each candidate business's own website and returns only the ones that match, each with the sentence and page that prove it. Businesses it cannot judge are shown as "couldn't tell", never guessed. The user sees a free match count first, pays only for matches, and can save the search to be alerted when new matches appear.

Who it's for. People who sell to local businesses and are not data engineers: AI automation agencies first, then web, SEO and local marketing agencies, then SaaS and B2B service companies selling to small businesses.

Why it wins. Every mainstream tool filters on map fields such as category, rating and review count. Google's own rules keep services and offers out of those fields, so relevance can only come from the website. Small Fish is built around that fact: relevance as the core, evidence on every row, honest uncertainty, and pricing aligned with what the buyer actually wants, which is matches.

The companion document. Market, competitor and customer research live in *Small Fish — PMM Foundations*. This document defines what to build, how it behaves, and how it makes money.

Vision and principles

In five years, Small Fish is the living map of what every local business actually does: the place anyone selling to a clinic, a salon or a contractor goes first, because it knows what each business offers, how it operates, and what just changed.

The arc

Stage	What Small Fish is	What the user feels
Launch	A search box that returns proven matches	"It found exactly who I sell to."
Year 1	Saved searches that bring new matches every week	"It keeps my pipeline full without me looking."
Year 2	A shared knowledge base of local businesses, read and refreshed continuously	"It already knows this market."
Year 3+	The relevance layer other tools plug into: CRMs, senders, AI agents	"Every tool I use knows who to talk to."

The compounding asset is the knowledge base. Every website Small Fish reads becomes structured facts (services offered, booking method, staff count, languages, specialties) that serve every future search. The more users search, the more the map fills in, and the cheaper and faster each search gets.

Principles

1. **Relevance over volume.** A short list the user trusts beats a long list they must re-check. We never compete on row count.
2. **Show the proof.** Every match carries the sentence and page that justify it. If we can't show it, we don't claim it.
3. **"Couldn't tell" is an honest answer.** A forced guess looks identical to a real match and destroys trust. Uncertainty is shown and never charged.
4. **Never invent.** No guessed emails, no made-up facts, no contact that the business did not publish.
5. **Charge for value, not effort.** Users pay for matches. Our scanning cost is our problem to engineer down, not theirs to fund.
6. **Plain English in, useful list out.** No filter trees, no tables to build, no prompt engineering.
7. **Respect the businesses we read.** Public business information only, polite crawling, and a way for any business to opt out.

Users and jobs

We build for the AI automation agency founder first; if the product works for someone who needs clinics with no online booking, it works for everyone who sells to local businesses.

Priority	User	Typical search	What "great" looks like
1	AI automation agency founder	"Dental clinics in Phoenix with no online booking"	A list they can pitch today, with the proof of the gap
2	Web, SEO and local marketing agency	"Roofers in Tampa with no service-area pages and no quote form"	Prospects whose site problem is their service
3	SaaS or B2B service selling to small businesses	"Independent vet clinics in Ohio offering exotic-pet care, not part of a group"	A qualified territory list, refreshed as it changes
4	Freelance list builder	Any client brief	A list the client never has to re-check

Jobs the product must nail

1. **Find:** return only businesses that match plain-English criteria in a chosen area.
2. **Prove:** show why each one matched, so the user trusts and can pitch it.
3. **Reach:** give the contact routes the business publishes: website, phone, email, contact form.
4. **Take it away:** export or push the list to where the user works.
5. **Keep it coming:** tell the user when new matches appear.

Who we are not building for

- Enterprise data teams wanting firmographic databases of large companies.
- Anyone seeking personal data about individuals, such as consumers or employees.
- Volume buyers who want 100,000 rows for mass email; they are better served by Outscraper.

The experience

One box, one count, one list: a new user goes from typing a search to holding a proven list in under three minutes, without creating an account until they want the list.

flowchart LR

```

A[Type a search] --> B[Confirm criteria]
B --> C[Free match count]
C --> D[Unlock matches]
D --> E[Review with proof]
E --> F[Export or push]
E --> G[Save search]
G --> H[Alerts on new matches]

```

The user never waits on a blank screen: the count and the first matches stream in as businesses are read.

1. Search

A single box on the home page, pre-filled with a rotating example. The user types naturally: *"Med spas in Dallas that offer Botox and don't have online booking."* Below the box, one-click templates by seller: "AI receptionist prospects", "Website redesign prospects", "Booking software prospects".

2. Confirm criteria

Small Fish turns the sentence into editable chips before spending anything:

- **Where:** Dallas, TX (city · 25 mi) — editable to a radius, county, state or drawn area
- **What:** Med spas — with suggested related types ("aesthetic clinics", "dermatology")
- **Must have:** offers Botox
- **Must not have:** online booking

Each criterion shows how it will be checked, in one line: "Looks for a booking widget, a Book Now link, or booking software on the site." Unclear criteria get a question, not a guess: "By *small*, do you mean fewer than 5 practitioners?"

3. Free match count

Small Fish scans a representative sample and shows an estimate within seconds, then firms it up as reading continues: **"About 140 matches out of 612 med spas. 38 couldn't be judged."** Three sample matches are shown in full, with their proof, for free. This is the

demo; no sign-up needed.

4. Unlock

One button: "Unlock 140 matches — uses 140 of your 300." New users sign up here with the free allowance. If the count exceeds their balance, they can unlock the top N, buy a pack, or upgrade.

5. Review with proof

A clean table. Each row: business name, area, match badge, the proof sentence with a link to the exact page, and published contacts. Clicking a row opens a side panel with every criterion checked, its evidence, and when the site was read. Filters on the table: matches, couldn't tell, has email, has phone, distance.

6. Take it away

CSV, Google Sheets, or push to HubSpot, Instantly, Smartlead, or a webhook. The export includes the proof columns, so a user's cold email can quote the evidence.

7. Save and get alerted

"Save this search" sets a weekly or monthly check. New matches arrive by email and in-app, and are charged only when unlocked. The saved-searches page shows each search's total matches, new this week, and last run.

Features

Ten features make up the complete product; the first six ship at launch, the rest follow once paying users ask for them.

#	Feature	Launch	Why it exists
1	Search builder	Yes	Plain English in, structured criteria out
2	Free match count	Yes	The hook and the demo
3	Results with proof	Yes	Trust; the core differentiator
4	Published contacts	Yes	Makes the list usable

#	Feature	Launch	Why it exists
5	Export	Yes	Gets the list into the user's workflow
6	Saved searches and alerts	Yes	Retention
7	Templates library	Soon after	Faster first value; SEO surface
8	Integrations (HubSpot, Instantly, Smartlead, Sheets)	Soon after	Removes the CSV step
9	Agency workspaces	Later	Multi-client use, the top price tier
10	API and MCP server	Later	Lets tools and AI agents use Small Fish

1. Search builder

- **Input:** free text, or a template.
- **Parsed into:** location (city, radius, county, state, drawn polygon), business types (primary plus suggested related types), must-have criteria, must-not-have criteria, optional preferences that rank but don't filter ("prefer multi-location").
- **Limits:** up to 5 criteria per search at launch; more criteria lower match rates and raise cost.
- **Check preview:** each criterion states what counts as proof, so users understand before paying.
- **Clarifying questions:** vague words ("small", "modern", "busy") trigger a single multiple-choice question.

2. Free match count

- Estimates from a sample of up to 60 businesses, shows a range, then tightens as reading continues.
- Shows three full sample matches with proof, free.
- Daily limit of free counts per user and per IP to control cost; logged-in free users get more.

3. Results with proof

- Every row: name, address, distance, match status, proof sentence, source page link, date read.
- Three statuses: **Match**, **No match** (hidden by default), **Couldn't tell** (shown, never charged).
- Side panel: each criterion, its verdict, its evidence, and its confidence in words ("clear", "likely"); never a raw number.
- Report a wrong result in one click; the credit is refunded automatically and the report feeds the quality system.

4. Published contacts

- Website, phone, public email, contact-form URL, and social profiles, each with where it was found.
- Only contacts the business published. No guessed emails, no personal emails of staff unless published by the business.
- Optional email verification before export, charged at cost.

5. Export

- CSV and Google Sheets at launch, including proof columns and a short "why it matched" line usable in outreach.
- Duplicate protection: businesses already exported by the user are flagged, not charged again.

6. Saved searches and alerts

- Weekly or monthly re-runs on the same criteria and area.
- Alerts show only new matches and businesses whose status changed ("added online booking — no longer matches").
- New matches are unlocked from the alert; nothing is charged until unlocked.

7. Templates library

Curated, editable searches for common sellers: AI receptionist prospects, booking software prospects, website redesign prospects, review management prospects, commercial HVAC buyers. Each template is also a public SEO page showing live counts by city.

8. Integrations

One-click push to HubSpot, Instantly, Smartlead, Google Sheets, and a generic webhook. Field mapping remembered per destination.

9. Agency workspaces

Separate client workspaces, shared credit pool, client-branded exports, and read-only client access to saved searches.

10. API and MCP server

The same search, count, and unlock operations over a REST API and an MCP server, so AI agents and automation tools (n8n, Make, Claude) can call Small Fish directly. Billed on the same per-match basis.

The engine

The engine reads each business once into a reusable fact profile, then judges every future search against that profile, so the second search in any market costs a fraction of the first. Cold reads cost about a cent per business; warm judgments about a fifth of a cent.

```
flowchart LR
  A[Parse search] --> B[Find candidates]
  B --> C{Profile fresh?}
  C -->|No| D[Fetch site pages]
  D --> E[Build fact profile]
  E --> F[Knowledge base]
  C -->|Yes| F
  F --> G[Judge criteria]
  G --> H[Match / No / Couldn't tell]
```

Every stage writes to the knowledge base, so work done for one user serves the next.

1. Parse the search

A small model turns the sentence into a structured query: location geometry, business types mapped to a category taxonomy, must-have and must-not-have criteria, and for each criterion a *check plan*: what evidence proves it, what evidence disproves it, and which pages to look at (services page, booking page, team page).

2. Find candidates

- **Primary:** Overture Maps and Foursquare open places data, loaded into our own database, filtered by geometry and category, deduplicated.
- **Gap-fill:** Google Places API Text Search for areas where open data looks thin; store place IDs only, as Google's terms require.
- **Websites:** taken from listings, then verified by fetching; missing sites found by a name-plus-city web search as a fallback.
- **Output:** a candidate set with a coverage estimate, so the count can say "612 med spas found in this area".

3. Read the site

- Fetch the homepage plus up to three linked pages chosen by the check plan (services, booking, about or team, contact).
- Render JavaScript only when the plain fetch returns an empty shell.
- Honour robots.txt, identify the crawler, throttle per domain.
- Detect known technology from page source: booking widgets (e.g. Calendly, Vagaro, Zocdoc, Mindbody, Jane, Boulevard), chat widgets, CMS, ad pixels. Technology detection proves many "has / has no" criteria without a model call.

4. Build the fact profile

One model pass per business extracts a structured profile: services and specialties, booking method, contact routes, staff or practitioner count, locations, languages, ownership hints (franchise, group, independent), and key quotes with their page URLs. The profile, not the raw page, is stored. It is reused for every future search until it goes stale.

5. Judge the criteria

- Cheap first: technology signals and keyword matches settle clear cases without a model call.
- Then a small model (Claude Haiku 4.5 class) judges the remaining criteria against the profile and quotes, returning a verdict, the proof quote, and a confidence word.
- A stronger model re-checks only the borderline cases and a random audit sample.

- **Absence criteria** ("no online booking") require positive proof of absence: the booking-relevant pages were read, no booking technology was found, and no booking link exists. If those pages couldn't be read, the verdict is "couldn't tell".

6. Freshness

- Profiles refresh on a 60-day default cycle, sooner for businesses in saved searches (weekly or monthly with the alert schedule).
- A cheap change check (page hash, headers) runs before a full re-read.
- Businesses marked closed in open data, or whose domain lapses, are retired.

7. Cost control

Lever	Effect
Shared knowledge base	Repeat searches in a market skip fetching and extraction
Tech detection and keywords before models	Settles a large share of criteria at near-zero cost
Small model by default, strong model on exceptions	Keeps per-business judging at ~\$0.001–0.002
Prompt caching and the batch API for alerts	Caching cuts repeated input to 10% of list price; batch halves all token costs (source)
Sampled counts	Free counts read a sample, not the full area
Scan budget per search	Stops runaway scanning on very rare criteria (see Pricing)

Cost per business, estimated

Step	Cold (first read)	Warm (profile cached)
Candidate data	~\$0 open data; ~\$0.002 with Google gap-fill	\$0
Fetching 2–4 pages	~\$0.0005	\$0

Step	Cold (first read)	Warm (profile cached)
Profile extraction (~6K tokens in, ~400 out, Haiku 4.5 at \$1/\$5 per million)	~\$0.008	\$0
Criteria judgment (~1.5K in, ~150 out)	~\$0.002	~\$0.002
Total	~\$0.01	~\$0.002

Model prices from [Anthropic pricing as reported July 2026](#). Treat these as planning estimates; measure the real figures in the first benchmark.

Edge cases

Every edge case resolves to one of three honest outcomes: a proven match, a clear no, or "couldn't tell" with the reason. The user is never charged for the last two.

Data and website problems

Situation	What Small Fish does	What the user sees
No website at all	Judges only listing-level criteria; anything needing the site → couldn't tell	"No website found." Optional toggle: <i>include no-website businesses</i> (valuable to web agencies)
Website is only a Facebook, Instagram or Linktree page	Reads the public page text where allowed; otherwise couldn't tell	"Social page only" tag
Site blocked by robots.txt	Does not fetch; couldn't tell	"Site doesn't allow reading"
Site down, parked or expired domain	Marks as likely closed or moved; retries once in 7 days	Excluded from matches; shown in couldn't tell with reason
JavaScript-only site	Renders it in a headless browser, at higher cost, capped per search	Normal result, slightly slower

Situation	What Small Fish does	What the user sees
Very thin site (one page, no services listed)	Judges what exists; missing evidence → couldn't tell, never no	"Too little on the site to judge"
Site in another language	Profile extraction handles major languages; proof quote shown in original with a translation	Proof in both languages
Listing and site disagree (address, hours, name)	Trusts the site for services, the listing for location; flags the mismatch	"Listing and website differ" note
Business closed	Retired when open data marks it closed or the site is gone	Never shown as a match

Business structure

Situation	What Small Fish does
Multi-location business, one website	One business, many locations; matches if any location falls in the area; lists the matching locations
Franchise or chain location	Detects franchise and group signals; the user can exclude chains in one click ("independent only")
Several businesses sharing one site (e.g. group practice)	Profiles the shared site once; flags "shared website" so claims aren't over-attributed
Duplicate listings for one business	Merged by domain, phone and address
Business moved outside the area	Uses the current address from the site if it disagrees with an old listing

Search problems

Situation	What Small Fish does
Vague criterion ("small", "modern site", "busy")	Asks one multiple-choice question before counting; saves the answer as the user's default
Criterion no website can prove ("revenue over \$1M", "owner is retiring")	Refuses politely at the confirm step and suggests a provable proxy ("has 5+ practitioners listed")
Absence criteria ("no online booking", "no chat widget")	Requires positive proof of absence from the relevant pages; otherwise couldn't tell
Very rare criteria (<3% match rate in sample)	Warns before unlocking: "Expect about 12 matches in this area"; suggests widening the area or loosening a criterion
Huge area (whole state, 20,000+ candidates)	Counts from a sample; unlocks run progressively, strongest matches first; user can stop at any point
Contradictory criteria	Flags the conflict at the confirm step
Search for people, not businesses ("find nurses", "homeowners")	Declined: Small Fish searches businesses only
Sensitive criteria about owners (religion, ethnicity, health, politics)	Declined: we judge what a business offers and how it operates, not who owns it

Billing and account problems

Situation	What Small Fish does
Match later proved wrong	One-click report; credit refunded automatically; case sent to quality review
Same business unlocked twice by the same user	Never charged twice within 12 months; shown as "already unlocked"
Business stops matching after unlock (e.g. added booking)	Saved-search alert tells the user; no charge involved

Situation	What Small Fish does
User runs out of matches mid-unlock	Unlocks stop at the balance; the rest wait, with a one-click top-up
Free-count abuse (bots, many accounts)	Rate limits per IP and device, CAPTCHA after repeated anonymous counts, cached counts for repeated queries
Reselling or bulk scraping via the API	Fair-use limits per plan; bulk resale needs the Agency plan

Business opt-out and data requests

Any business can ask to be excluded via a public form. Small Fish verifies ownership through the listed domain or phone, stops reading the site, removes its profile, and suppresses it from all searches within 7 days.

Quality system

Precision is the product: at least 90% of shown matches must be correct when checked by hand, and no change ships if it lowers that number.

Targets

Measure	Definition	Launch target	Year 1 target
Match precision	Share of shown matches a human agrees with	≥ 90%	≥ 95%
Couldn't-tell rate	Share of candidates we decline to judge	≤ 25%	≤ 15%
False "no" rate	True matches we wrongly reject (sampled)	≤ 15%	≤ 10%
Proof validity	Proof quote exists verbatim on the linked page	100%	100%
Wrong-result reports	Reports per 1,000 unlocked matches	≤ 20	≤ 10

Precision matters more than recall: a missed match costs the user nothing, a wrong match costs them a bad email and their trust.

The benchmark set

- **Launch set:** 10 common searches (e.g. med spas with no online booking, dentists offering implants, HVAC doing commercial work, independent vet clinics), 100 businesses each, labelled by hand.
- **Growth:** every confirmed user report becomes a labelled case.
- **Published:** precision and couldn't-tell rates for the public benchmark searches appear on the website, updated monthly.

Gates and monitoring

1. **Regression gate:** every change to prompts, models or extraction runs against the benchmark; a precision drop of more than 1 point blocks release.
2. **Proof check:** every proof quote is verified to exist on its page before display; failures become couldn't tell.
3. **Audit sample:** 1% of live matches are re-judged daily by a stronger model; disagreement above 5% for any criterion type alerts the team.
4. **Drift watch:** couldn't-tell and report rates tracked per criterion type, per category and per region.

Learning loop

Reports and refunds feed three places: the benchmark, the criterion check plans ("booking link hidden in footer" becomes a rule), and the technology detector's catalogue of booking and chat tools.

Pricing

One unit, the **match**, sold in four simple plans from \$0 to \$199 a month, plus no-subscription packs. Non-matches and "couldn't tell" are always free. Target blended gross margin: 70% or better.

The unit

A **match** is one business the user unlocks that meets their criteria, with proof. Counting, previews, re-viewing, re-exporting and alerts are free. The user pays for exactly what they came for, which is the core promise against tools that bill per row.

Plans

Plan	Monthly	Annual (per month)	Matches / month	Effective price per match	Saved searches	Includes
Free	\$0	\$0	25	—	1, monthly alerts	Free counts (5/day), CSV
Starter	\$29	\$24	300	\$0.097	5, weekly alerts	CSV, Google Sheets, email support
Growth	\$79	\$66	1,000	\$0.079	25, weekly alerts	HubSpot, Instantly, Smartlead, webhook; templates library
Agency	\$199	\$166	3,000	\$0.066	Unlimited	10 client workspaces, white-label exports, API and MCP, priority support

Packs, no subscription: 100 matches for \$15, 500 for \$60, valid 12 months. They answer the buyer who needs "only a few lists" and resents a subscription, while pricing per match higher than plans to make subscribing the obvious deal.

Extra matches on a plan: billed at the plan's effective rate, with optional auto top-up. Unused matches roll over one month, capped at one month's allowance.

Rules that protect margin, stated openly

1. **Rare-match searches cost 2 matches per result.** If the free count's sample shows a match rate below 5%, the confirm screen says so before unlocking: "This is a rare search: each result uses 2 matches." Rare searches cost us 4× more to find.
2. **Scan budget.** An unlock scans at most 20 candidates per requested match, then stops and reports how many it found.
3. **Email verification at cost:** about half a cent per address, shown before export.
4. **Fair use on free counts:** 5 a day anonymous, 20 a day logged in, cached per query for 7 days.

Unit economics

Cost per match equals cost per business judged divided by the match rate. Cold reads cost ~\$0.01 per business and warm ones ~\$0.002 (see The engine). "Blended" assumes half of reads hit a fresh cached profile, which should hold by month 6 in our core niches.

Match rate	Cold cost per match	Warm cost per match	Blended cost per match	Gross margin at Growth (\$0.079)
30% (broad search)	\$0.033	\$0.007	\$0.020	75%
15% (typical)	\$0.067	\$0.013	\$0.040	49%
5% (rare, charged 2×)	\$0.200	\$0.040	\$0.120	24% on \$0.158

Two things lift real margin above this table. First, **breakage**: subscription tools typically see a meaningful share of allowances go unused; at 30% unused, the typical-case margin rises to about 65%. Second, **knowledge-base growth**: every month more reads are warm, pushing the typical case toward 80%. The rare-search surcharge keeps the worst case positive.

Kill threshold: if measured blended cost per match exceeds \$0.05 after month 3, raise prices or reduce the Growth allowance before scaling acquisition.

Against the competition

Option	Headline price	Cost per <i>useful</i> lead	Hidden cost
Outscraper	~\$3 per 1,000 rows	~\$0.02 if 15% are relevant	Hours of manual website checking
Scrap.io Basic	\$35/month for 10,000 exports	~\$0.02 if 15% are relevant	Same manual checking; no proof
Exa Websets Core	\$49/month ≈ 800 verified results	~\$0.06	Not local-native; 100-result cap per Webset
Freelance list builder	\$13–45 per hour	~\$2 per qualified lead at 10 an hour	Slow; inconsistent

Option	Headline price	Cost per <i>useful</i> lead	Hidden cost
Small Fish Growth	\$79/month for 1,000 matches	\$0.079, proven	None: no checking needed

The honest pitch: raw rows are cheaper, but Small Fish is the cheapest way to get a lead you don't have to check. Against Exa, the price is close and the product is built for local search. Starter at \$29 undercuts Scrap.io's \$35 entry.

Revenue model at 700 paying users

Assuming a mix of 50% Starter, 35% Growth and 15% Agency, average revenue per paying user is about \$72 a month: roughly **\$50K MRR** at 700 users, before packs and extra matches.

Pricing experiments to run after launch

- Starter at \$29 vs \$39.
- Free plan at 25 vs 10 matches (activation vs cost).
- Rare-search surcharge at 2× vs a flat "precision search" add-on.

Metrics

The north star is **proven matches exported per week**: it rises only when users find relevant businesses, trust them, and take them into their work.

Area	Metric	Target at 6 months
North star	Proven matches exported per week	25,000
Activation	Visitors who run a free count	30% of home-page visitors
Activation	Free count → sign-up	20%
Activation	Sign-up → first export within 24 hours	50%
Conversion	Free → paid within 30 days	6%
Quality	Match precision on the benchmark	≥ 90%
Quality	Wrong-result reports per 1,000 matches	≤ 20

Area	Metric	Target at 6 months
Retention	Paid users with an active saved search	60%
Retention	Monthly paid churn	≤ 6%
Economics	Blended cost per match	≤ \$0.04
Economics	Share of reads served warm from the knowledge base	≥ 50%
Economics	Gross margin after breakage	≥ 65%, rising to 75%

Roadmap

Five stages, each unlocked by a gate, not a date: prove accuracy, launch the core, grow the knowledge base, then become the relevance layer other tools use.

Stage	Build	Gate to move on
0. Prove it (weeks 1–2)	Benchmark: 3 niche searches through Small Fish, Scrap.io and Exa Websets; measure cost per business	Small Fish precision ≥ 90% and clearly ahead on true matches per search
1. Launch (weeks 3–8)	Features 1–6: search builder, free count, results with proof, contacts, CSV and Sheets export, saved searches; Free, Starter, Growth plans; packs; 20 template pages	50 paying users; precision holds ≥ 90% live
2. Grow (months 3–6)	Templates library, integrations (HubSpot, Instantly, Smartlead, webhook), Agency plan and workspaces, published benchmark page, UK and Canada coverage	300 paying users; churn ≤ 6%; blended cost per match ≤ \$0.04

Stage	Build	Gate to move on
3. Compound (months 6-18)	API and MCP server, change alerts ("added booking", "new location"), knowledge base covering top 50 US metros in core niches, Australia	1,000 paying users; $\geq 50\%$ warm reads
4. The relevance layer (18 months+)	Small Fish inside other tools: CRM and sender integrations that call relevance on demand, partner marketplace listings, data licensing of derived facts	The dream: "who should I sell to" answered wherever a seller works

What we will not build early

Email sending, a CRM, contact databases of individuals, and volume scraping features. Each would pull the product away from relevance and toward commodity competition.

Non-goals, risks and open questions

Small Fish stays narrow on purpose: relevance for local businesses, proven and priced per match.

Non-goals

- Not a sender, a CRM or a sequencer; we export to them.
- Not a people database; we do not search or sell data about individuals.
- Not a volume scraper; we don't compete on rows per dollar.
- Not a firmographic database for large companies.

Top risks

Risk	Early warning	Response
Accuracy isn't clearly better than Exa or Scrap.io filters	Stage 0 benchmark	Don't launch; narrow to the niches where we win

Risk	Early warning	Response
Open data misses too many small businesses	Coverage below 70% of Google's count in a niche	Increase Google Places gap-fill for that niche; revisit costs
Blended cost per match above \$0.05	Monthly cost report	Raise prices or trim allowances before scaling
Scrap.io or Exa ships local AI criteria	Competitor changelogs	Lean on proof UX, honesty and pay-per-match; publish benchmarks
Low-price churn	Monthly churn above 8%	Push saved searches and annual plans; lift the Agency tier

Open questions

- ☐ What precision do we achieve on the first 3 benchmark searches, and at what cost per business?
- ☐ How does open-data coverage compare with Google Maps in our first 3 niches?
- ☐ Do buyers accept the rare-search 2× rule when it's shown up front?
- ☐ Is 25 free matches the right activation hook, or does it attract list farmers?
- ☐ Which 5 templates drive the most first searches?